SML/NJ Language Processing Tools:
User Guide

Aaron Turon
adrassi@gmail.com

John Reppy
jhr@cs.uchciago.edu

November 2013 — 110.76+ release

ii

Contents

Chapter 1

Overview

In software, language recognition is ubiquitous: nearly every program deals at some level with structured input given in textual form. The simplest recognition problems can be solved directly, but as the complexity of the language grows, recognition and processing become more difficult.

Although sophisticated language processing is sometimes done by hand, the use of scanner and parser generators[1] is more common. The Unix tools `lex` and `yacc` are the archetypical examples of such generators. Tradition has it that when a new programming language is introduced, new scanner and parser generators are written in that language, and generate code for that language. Traditional *also* has it that the new tools are modeled after the old `lex` and `yacc` tools, both in terms of the algorithms used, and often the syntax as well. The language Standard ML is no exception: `ml-lex` and `ml-yacc` are the SML incarnations of the old Unix tools.

This manual describes two new tools, `ml-ulex` and `ml-antlr`, that follow tradition in separating scanning from parsing, but break from tradition in their implementation: `ml-ulex` is based on *regular expression derivatives* rather than subset-construction, and `ml-antlr` is based on $LL(k)$ parsing rather than $LALR(1)$ parsing.

Motivation

Most parser generators use some variation on *LR* parsing, a form of *bottom-up* parsing that tracks possible interpretations (reductions) of an input phrase until only a single reduction is possible. While this is a powerful technique, it has the following downsides:

- Compared to predictive parsing, it is more complicated and difficult to understand. This is particularly troublesome when debugging an *LR*-ambiguous grammar.

- Because reductions take place as late as possible, the choice of reduction cannot

---

[1]"Scanner generator" and "parser generator" will often be shortened to "scanner" and "parser" respectively. This is justified by viewing a parser generator as a parameterized parser.

depend on any semantic information; such information would only become available *after* the choice was made.

- Similarly, information flow in the parser is strictly bottom-up. For (syntactic or semantic) context to influence a semantic action, higher-order programming is necessary.

The main alternative to *LR* parsing is the top-down, *LL* approach, which is commonly used for hand-coded parsers. An *LL* parser, when faced with a decision point in the grammar, utilizes lookahead to unambiguously predict the correct interpretation of the input. As a result, *LL* parsers do not suffer from the problems above. *LL* parsers have been considered impractical because the size of their prediction table is exponential in k — the number of tokens to look ahead — and many languages need $k > 1$. However, Parr showed that an approximate form of lookahead, using tables linear in k, is usually sufficient.

To date, the only mature *LL* parser based on Parr's technique is his own parser, `antlr`. While `antlr` is sophisticated and robust, it is designed for and best used within imperative languages. The primary motivation for the tools this manual describes is to bring practical *LL* parsing to a functional language. Our hope with `ml-ulex` and `ml-antlr` is to modernize and improve the Standard ML language processing infrastructure, while demonstrating the effectiveness of regular expression derivatives and $LL(k)$ parsing. The tools are more powerful than their predecessors, and they raise the level of discourse in language processing.

Chapter 2

ML-ULex

Lexers analyze the lexical structure of an input string, and are usually specified using regular expressions. ML-ULEX is a lexer generator for Standard ML. The module it generates will contain a type `strm` and a function

```
val lex : AntlrStreamPos.sourcemap -> strm
        -> lex_result * AntlrStreamPos.span * strm
```

where `lex_result` is a type that must be defined by the user of `ml-ulex`. Note that the lexer always returns a token: we assume that end-of-file will be explicitly represented by a token. Compared to ML-Lex, `ml-ulex` offers the following improvements:

- Unicode is supported under the UTF8 encoding.

- Intersection and negation of REs are supported.

- Token position spans are automatically generated.

- The specification format is somewhat cleaner.

- The code base is much cleaner, and supports multiple back-ends, including DFA graph visualization and interactive testing of rules.

The tool is invoked from the command-line as follows:

```
ml-ulex [options] file
```

where `file` is the name of the input `ml-ulex` specification, and where `options` may be any combination of:

`--dot`	generate DOT output (for graphviz; see `http://www.graphviz.org`). The produced file will be named `file.dot`, where `file` is the input file.
`--match`	enter interactive matching mode. This will allow interactive testing of the machine; presently, only the `INITIAL` start state is available for testing (see Section 2.2.6 for details on start states).
`--ml-lex-mode`	operate in `ml-lex` compatibility mode. See Section 2.5 for details.
`--table-based`	generate a table-based lexer.
`--fn-based`	generate a lexer that represents states as functions and transitions as tail calls.
`--minimize`	generate a minimal machine. Note that this is slow, and is almost never necessary.
`--strict-sml`	generate strict SML (*i.e.*, do not use SML/NJ extensions). This flag is useful if you want to use the output with a different SML system.

The output file will be called `file.sml`.

2.1   Specification format

A `ml-ulex` specification is a list of semicolon-terminated *declarations*. Each declaration is either a *directive* or a *rule*. Directives are used to alter global specification properties (such as the name of the module that will be generated) or to define named regular expressions. Rules specify the actual reguluar expressions to be matched. The grammar is given in Figure 2.1.

There are a few lexical details of the specification format worth mentioning. First, SML-style comments (`(* ... *)`) are treated as ignored whitespace anywhere they occur in the specification, *except* in segments of code. The *ID* symbol used in the grammar stands for alpha-numeric-underscore identifiers, starting with an alpha character. The *code* symbol represents a segment of SML code, enclosed in parentheses. Extra parentheses occuring within strings or comments in code need not be balanced.

A complete example specification appears in Chapter 5.

$$
\begin{array}{rcl}
\textit{spec} & ::= & (\ \textit{declaration}\ ;\ )^* \\
\textit{declaration} & ::= & \textit{directive} \\
& | & \textit{rule} \\
\textit{directive} & ::= & \texttt{\%arg}\ \textit{code} \\
& | & \texttt{\%defs}\ \textit{code} \\
& | & \texttt{\%let}\ \textit{ID} = \textit{re} \\
& | & \texttt{\%name}\ \textit{ID} \\
& | & \texttt{\%header}\ \textit{code} \\
& | & \texttt{\%states}\ \textit{ID}^{+} \\
\textit{code} & ::= & (\ \ldots\ ) \\
\textit{rule} & ::= & (\texttt{<}\ \textit{ID}\ (\ ,\ \textit{ID}\ )^*\ \texttt{>})^?\ \textit{re}\ \texttt{=>}\ \textit{code} \\
& | & (\texttt{<}\ \textit{ID}\ (\ ,\ \textit{ID}\ )^*\ \texttt{>})^?\ \texttt{<<EOF>>}\ \texttt{=>}\ \textit{code} \\
\textit{re} & ::= & \textit{CHAR} \\
& | & \texttt{"}\ \textit{SCHAR}^*\ \texttt{"} \\
& | & (\ \textit{re}\ ) \\
& | & \texttt{[}\ (\ \texttt{-}\ |\ \texttt{\textasciicircum}\ )^?\ (\ \textit{CCHAR}\ \texttt{-}\ \textit{CCHAR}\ |\ \textit{CCHAR}\ )^{+}\ \texttt{-}^?\ \texttt{]}
\end{array}
$$

		a character class
	$\{\ \textit{ID}\ \}$	`%let`-bound RE
	.	wildcard (any single character including `\n`)
	re *	Kleene-closure (0 or more)
	re ?	optional (0 or 1)
	re +	positive-closure (1 or more)
	re { *NUM* }	exactly *NUM* repetitions
	re { NUM_1, NUM_2 }	between NUM_1 and NUM_2 repetitions
	re re	concatenation
	$\sim$ *re*	negation
	re & *re*	intersection
	re | *re*	union

$$
\begin{array}{rcl}
\textit{CHAR} & ::= & \textit{any printable character not one of}\ \ \texttt{\textasciicircum\ <\ >\ \textbackslash\ (\ )\ \{\ \}\ [\ \&\ |\ *\ ?} \\
& & \hspace{4em} \texttt{+\ "\ .\ ;\ =\ \textasciitilde} \\
& | & \textit{an SML or Unicode escape code} \\
\textit{CCHAR} & ::= & \textit{any printable character not one of}\ \texttt{\textasciicircum\ -\ ]\ \textbackslash} \\
& | & \textit{an SML or Unicode escape code} \\
\textit{SCHAR} & ::= & \textit{any printable character not one of}\ \texttt{"\ \textbackslash} \\
& | & \textit{an SML or Unicode escape code} \\
\textit{NUM} & ::= & \textit{one or more digits}
\end{array}
$$

Figure 2.1: The `ml-ulex` grammar

2.2   Directives

2.2.1   The %arg directive

Specifies an additional curried parameter, appearing after the sourcemap parameter, that will be passed into the `lex` function and made available to all lexer actions.

2.2.2   The %defs directive

The `%defs` directive is used to include a segment of code in the generated lexer module, as in the following example:

```
%defs (
  type lex_result = CalcParseToks.token
  fun eof() = CalcParseTokens.EOF
  fun helperFn x = (* ... *)
)
```

The definitions must at least fulfill the following signature:

```
type lex_result
val eof : unit -> lex_result
```

unless EOF rules are specified, in which case only the lex_result type is needed (see Section 2.3.2). All semantic actions must yield values of type lex_result. The `eof` function is called by `ml-ulex` when the end of file is reached – it acts as the semantic action for the empty input string. All definitions given will be in scope for the rule actions (see Section 2.3).

2.2.3   The %let directive

Use `%let` to define named abbreviations for regular expressions; once bound, an abbreviation can be used in further `%let`-bindings or in rules. For example,

```
%let digit = [0-9];
```

introduces an abbreviation for a regular expression matching a single digit. To use abbreviations, enclose their name in curly braces. For example, an additional `%let` definition can be given in terms of `digit`,

```
%let int = {digit}+;
```

which matches arbitrary-length integers. Note that scoping of let-bindings follows standard SML rules, so that the definition of `int` must appear after the definition of `digit`.

2.2.4   The %name directive

The name to use for the generated lexer module is specified using `%name`.

2.2.5 The %header directive

The %header directive allows one to specify the lexer's header. This directive is useful when you want to use a functor for the lexer. For example,

```
%header (
  functor ExampleLexFn (Extras : EXTRA_SIG));
```

Note that it is an error to include both the %header and %name directive in the same lexer specification.

2.2.6 The %states directive

It is often helpful for a lexer to have multiple *start states*, which influence the regular expressions that the lexer will match. For instance, after seeing a double-quote, the lexer might switch into a STRING start state, which contains only the rules necessary for matching strings, and which returns to the standard start state after the closing quote.

Start states are introduced via %states, and are named using standard identifiers. There is always an implicit, default start state called INITIAL. Within a rule action, the function YYBEGIN can be applied to the name of a start state to switch the lexer into that state; see 2.3.3 for details on rule actions.

2.3 Rules

In general, when lex is applied to an input stream, it will attempt to match a prefix of the input with a regular expression given in one of the rules. When a rule is matched, its *action* (associated code) is evaluated and the result is returned. Hence, all actions must belong to the same type.Rules are specified by an optional list of start states, a regular expression, and the action code. The rule is said to "belong" to the start states it lists. If no start states are specified, the rule belongs to *all* defined start states.

Rule matching is determined by three factors: start state, match length, and rule order. A rule is only considered for matching if it belongs to the lexer's current start state. If multiple rules match an input prefix, the rule matching the longest prefix is selected. In the case of a tie, the rule appearing first in the specification is selected.

For example, suppose the start state FOO is defined, and the following rules appear, with no other rules belonging to FOO:

```
<FOO> a+    => ( Tokens.as );
<FOO> a+b+  => ( Tokens.asbs );
<FOO> a+bb* => ( Tokens.asbs );
```

If the current start state is not FOO, none of the rules will be considered. Otherwise, on input "aabbbc" all three rules are possible matches. The first rule is discarded, since the others match a longer prefix. The second rule is then selected, because it matches the same prefix as the third rule, but appears earlier in the specification.

$$\llbracket . \rrbracket \ = \ \Sigma \qquad\qquad \llbracket re * \rrbracket \ = \ \bigcup_{i=0}^{\infty} \llbracket re \rrbracket^{i}$$

$$\llbracket re_1 \ re_2 \rrbracket \ = \ \llbracket re_1 \rrbracket \cdot \llbracket re_2 \rrbracket$$

$$\llbracket \sim re \rrbracket \ = \ \left(\bigcup_{i=0}^{\infty} \Sigma^{i} \right) \setminus \llbracket re \rrbracket \qquad\qquad \llbracket re + \rrbracket \ = \ \bigcup_{i=1}^{\infty} \llbracket re \rrbracket^{i}$$

$$\llbracket re_1 \ \& \ re_2 \rrbracket \ = \ \llbracket re_1 \rrbracket \cap \llbracket re_2 \rrbracket \qquad\qquad \llbracket re \ \{n\} \rrbracket \ = \ \llbracket re \rrbracket^{n}$$

$$\llbracket re_1 \ | \ re_2 \rrbracket \ = \ \llbracket re_1 \rrbracket \cup \llbracket re_2 \rrbracket \qquad\qquad \llbracket re \ \{n, \ m\} \rrbracket \ = \ \bigcup_{i=n}^{m} \llbracket re \rrbracket^{i}$$

$$\llbracket re \ ? \rrbracket \ = \ \llbracket re \rrbracket \cup \{\epsilon\}$$

Figure 2.2: Semantics for regular expressions

2.3.1 Regular expression syntax

The syntax of regular expressions is given in Figure 2.1; constructs are listed in precedence order, from most tightly-binding to least. Escape codes are the same as in SML, but also include \uxxxx and \Uxxxxxxxx, where xxxx represents a hexidecimal number which in turn represents a Unicode symbol. The specification format itself freely accepts Unicode characters, and they may be used within a quoted string, or by themselves.

The semantics for `ml-ulex` regular expressions are shown in Figure 2.2; they are standard. Some examples:

$$
\begin{aligned}
\texttt{0|1|2|3} \quad &\textit{denotes} \quad \{0, 1, 2, 3\} \\
\texttt{[0123]} \quad &\textit{denotes} \quad \{0, 1, 2, 3\} \\
\texttt{0123} \quad &\textit{denotes} \quad \{0123\} \\
\texttt{0*} \quad &\textit{denotes} \quad \{\epsilon, 0, 00, \dots\} \\
\texttt{00*} \quad &\textit{denotes} \quad \{0, 00, \dots\} \\
\texttt{0+} \quad &\textit{denotes} \quad \{0, 00, \dots\} \\
\texttt{[0-9]\{3\}} \quad &\textit{denotes} \quad \{000, 001, 002, \dots, 999\} \\
\texttt{0*\&(..)*} \quad &\textit{denotes} \quad \{\epsilon, 00, 0000, \dots\} \\
\texttt{\^{}(abc)} \quad &\textit{denotes} \quad \Sigma^{*} \setminus \{abc\} \\
\texttt{[\^{}abc]} \quad &\textit{denotes} \quad \Sigma \setminus \{a, b, c\}
\end{aligned}
$$

2.3.2 EOF rules

It is sometimes useful for the behavior of a lexer when it reaches the end-of-file to change depending on the current start state. Normally, there is a single user-defined `eof` function that defines EOF behavior, but EOF rules can be used to be more selective, as in the following example:

```
<INITIAL> <<EOF>> => ( Tok.EOF );
<COMMENT> <<EOF>> => ( print "Error: unclosed comment";
                       Tok.EOF );
```

Other than the special <<EOF>> symbol, EOF rules work exactly like normal rules.

2.3.3 Actions

Actions are arbitrary SML code enclosed in parentheses. The following names are in scope:

`YYBEGIN`	a function taking a start state and returning unit; changes to that start state.
`yysetStrm`	a function taking a `ULexBuffer.stream` and returning unit; changes the current input source. The functions `yystreamify`, `yystreamifyInstream` and `yystreamifyReader` can be used to construct the stream; they work identically to the corresponding functions described in Section 2.4
`yytext`	the matched text as a string.
`yysubstr`	the matched text as a substring (avoids copying).
`yyunicode`	the matched Unicode text as a list of `Word.word` values
`continue`	a unit to `lex_result` function which recursively calls the lexer on the input following the matched prefix, and returns its result. The span for the resulting token begins at the left border of the match that calls `continue`,
`skip`	identical to `continue`, but moves forward the left marker for the position span of the returned token. Thus `skip` should be used for skipping whitespace.
`yysm`	the sourcemap for the lexer, to be used with the functions in the `AntlrSourcePos` module.
`yypos`	the position of the left border of the matched RE, starting from 0.
`yylineno`	the current line number, starting from 1.
`yycolno`	the current column number, starting from 1.
`?`	any name bound in the `%defs` section.

2.4 Using the generated code

The generated lexer module has a signature including the following:

```
type prestrm
type strm = prestrm * start_state

val streamify        : (unit -> string) -> strm
val streamifyReader  : (char, 'a) StringCvt.reader -> 'a -> strm
val streamifyInstream : TextIO.instream -> strm

val lex     : AntlrStreamPos.sourcemap -> strm ->
```

```
lex_result * AntlrStreamPos.span * strm
```

where `lex_result` is the result type of the lexer actions, and `start_state` is an algebraic datatype with nullary constructors for each defined start state. Note that `lex_result` must be defined as a type using the `%defs` directive. In this interface, lexer start states are conceptually part of the input stream; thus, from an external viewpoint start states can be ignored. However, it is sometimes helpful to control the lexer start state externally, allowing contextual information to influence the lexer. This is why the `strm` type includes a concrete `start_state` component.

Note that the `AntlrStreamPos` module is part of the `ml-lpt-lib` library described in Chapter 4. An `AntlrStreamPos.sourcemap` value, combined with an `AntlrStreamPos.pos` value, compactly represents position information (line number, column number, and so on). An `AntlrStreamPos.span` is a pair of `pos` values.

2.5   `ml-lex` compatibility

Running `ml-ulex` with the `--ml-lex-mode` option will cause it to process its input file using the ML-Lex format, and interpret the actions in a ML-Lex-compatible way. The compatibility extends to the bugs in ML-Lex, so in particular `yylineno` starts at 2 in `--ml-lex-mode`.

Chapter 3

ML-Antlr

Parsers analyze the syntactic structure of an input string, and are usually specified with some variant of context-free grammars. `ml-antlr` is a parser generator for Standard ML based on Terence Parr's variant of $LL(k)$ parsing. The details of the parsing algorithm are given in the companion implementation notes; the practical restrictions on grammars are discussed in Section 3.5. A parser generated by `ml-antlr` is a functor; it requires a module with the `ANTLR_LEXER` signature:

```
signature ANTLR_LEXER = sig
  type strm
  val getPos : strm -> AntlrStreamPos.pos
end
```

Applying the parser functor will yield a module containing a `parse` function:

```
val parse :
  (Lex.strm -> ParserToks.token * AntlrStreamPos.span * Lex.strm) ->
  Lex.strm ->
  result_ty option * strm * ParserToks.token AntlrRepair.repair list
```

where `result_ty` is determined by the semantic actions for the parser. The `ParserTokens` module is generated by `ml-antlr` (see Section 3.7) and the `AntlrRepair` module is available in the `ml-lpt` library (see Chapter 4).

Notable features of `ml-antlr` include:

- Extended BNF format, including Kleene-closure (*), positive closure (+), and optional (?) operators.

- Robust, automatic error repair.

- Selective backtracking.

- "Inherited attributes": information can flow downward as well as upward during a parse.

- Semantic predicates: a syntactic match can be qualified by a semantic condition.

11

- Grammar inheritence.

- Convenient default actions, especially for EBNF constructions.

- Convenient abbreviations for token names (*e.g.,* `"("` rather than `LP`)

The tool is invoked from the command-line as follows:

```
ml-antlr [options] file
```

where `file` is the name of the input `ml-ulex` specification, and where `options` may be any combination of:

`--dot`	generate DOT output (for graphviz; see `http://www.graphviz.org`). The produced file will be named `file.dot`, where `file` is the input file.
`--latex`	generate a simple LATEXversion of the grammar, named `file.tex`.
`--unit-actions`	ignore the action code in the grammar, and instead return `()` for every production.

The output file will be called `file.sml`.

3.1   Background definitions

Before describing `ml-antlr`, we need some terminology. A *context-free grammar* (CFG) is a set of *token* (or *terminal*) symbols, a set of *nonterminal* symbols, a set of *productions*, and a start symbol S, which must be a nonterminal. The general term *symbol* refers to both tokens and nonterminals. A production relates a nonterminal A to a string of symbols α; we write this relation as $A \rightarrow \alpha$. Suppose $\alpha A \beta$ is a symbol string, and A is a nonterminal symbol. We write $\alpha A \beta \Rightarrow \alpha \gamma \beta$ if $A \rightarrow \gamma$ is a production; this is called a one-step derivation. In general, a CFG generates a language, which is a set of token strings. The strings included in this language are exactly those token string derived in one or more steps from the start symbol S.

A parser recognizes whether an input string is in the language generated by a given CFG, usually computing some value (such as a parse tree) while doing so. The computations performed during a parse are called *semantic actions* (or just *actions*).

3.2   Specification format

A `ml-antlr` specification is a list of semicolon-terminated *declarations*. Each declaration is either a *directive* or a *nonterminal definition*. Directives are used to alter global specification properties (such as the name of the functor that will be generated) or to

$$
\begin{array}{rcl}
\textit{spec} & ::= & (\textit{ declaration } ; \textit{ })^* \\
\textit{declaration} & ::= & \textit{directive} \\
& | & \textit{nonterminal} \\
\textit{directive} & ::= & \texttt{\%defs } \textit{code} \\
& | & \texttt{\%entry } \textit{ID } (\textit{ , ID })^* \\
& | & \texttt{\%import } \textit{STRING } (\texttt{ \%dropping } \textit{symbol}^+ \textit{ })^? \\
& | & \texttt{\%keywords } \textit{symbol } (\textit{ , symbol })^* \\
& | & \texttt{\%name } \textit{ID} \\
& | & \texttt{\%header } \textit{code} \\
& | & \texttt{\%refcell } \textit{ID } : \textit{ monotype } = \textit{code} \\
& | & \texttt{\%start } \textit{ID} \\
& | & \texttt{\%tokens } : \textit{ tokdef } (\textit{ | tokdef })^* \\
& | & \texttt{\%nonterms } : \textit{ datacon } (\textit{ | datacon })^* \\
\textit{code} & ::= & (\ldots) \\
\textit{tokdef} & ::= & \textit{datacon } (\textit{ (STRING) })^? \\
\textit{datacon} & ::= & \textit{ID} \\
& | & \textit{ID } \text{of } \textit{monotype} \\
\textit{monotype} & ::= & \text{standard SML syntax for monomorphic types} \\
\textit{nonterminal} & ::= & \textit{ID formals}^? : \textit{ prodlist} \\
\textit{formals} & ::= & (\textit{ ID } (\textit{ , ID })^*) \\
\textit{prodlist} & ::= & \textit{production } (\textit{ | production })^* \\
\textit{production} & ::= & \texttt{\%try}^? \textit{ named-item}^* (\texttt{ \%where } \textit{code })^? (\texttt{ => } \textit{code })^? \\
\textit{named-item} & ::= & (\textit{ ID } :)^? \textit{ item} \\
\textit{item} & ::= & \textit{prim-item } \texttt{?} \\
& | & \textit{prim-item } \texttt{+} \\
& | & \textit{prim-item } \texttt{*} \\
\textit{prim-item} & ::= & \textit{symbol } (\texttt{ @ } \textit{code })^? \\
& | & (\textit{prodlist }) \\
\textit{symbol} & ::= & \textit{ID} \\
& | & \textit{STRING} \\
\textit{ID} & ::= & \text{standard SML identifier} \\
\textit{STRING} & ::= & \text{standard SML double-quoted string}
\end{array}
$$

Figure 3.1: The `ml-antlr` grammar

define supporting infrastructure for the grammar. The nonterminal definitions specify the grammar itself. The grammar for `ml-antlr` is given in Figure 3.1.

SML-style comments ((* ... *)) are treated as ignored whitespace anywhere they occur in the specification, *except* in segments of code. The *code* symbol represents a segment of SML code, enclosed in parentheses. Extra parentheses occuring within strings or comments in code need not be balanced. A complete example specification appears in Chapter 5.

Most `ml-antlr` declarations are *cumulative*: they may appear multiple times in a grammar specification, with each new declaration adding to the effect of the previous ones. Thus, for instance, the specification fragment

```
%tokens : foo ;
%tokens : bar of string ;
```

is equivalent to the single directive

```
%tokens : foo | bar of string ;
```

and similarly for nonterminal definitions and so on. All declarations are cumulative except for the `%start` and `%name` directives. The reason for treating specifications in this way is to give the `%import` directive very simple semantics, as described below.

3.3   Directives

3.3.1   The `%defs` directive

The `%defs` directive is used to include a segment of code in the generated parser:

```
%defs (
  fun helperFn x = (* ... *)
);
```

All definitions given will be in scope for the semantic actions (see Section 3.4.5).

3.3.2   The `%entry` directive

It is often useful to parse input based on some fragment of a grammar. When a nonterminal is declared to be an *entry point* for the grammar via `%entry`, `ml-antlr` will generate a separate `parse` function that expects the input to be a string derived from that nonterminal. Given a grammar with a nonterminal exp and the declaration

```
%entry exp;
```

the generated parser will include a function

```
val parseexp :
  (Lex.strm -> ParserToks.token * AntlrStreamPos.span * Lex.strm) ->
  Lex.strm ->
  exp_ty option * strm * ParserToks.token AntlrRepair.repair list
```

where `exp_ty` is the type of the actions for the `exp` nonterminal. Note that if `exp` has inherited attributes (Section 3.4.2) they will appear as a tuple argument, curried after the lexer argument:

```
val parseexp :
  (Lex.strm -> ParserToks.token * AntlrStreamPos.span * Lex.strm) ->
  attributes ->
  Lex.strm ->
  exp_ty option * strm * ParserToks.token AntlrRepair.repair list
```

Finally, the *start* symbol (Section 3.3.8) is always an entry point to the grammar, but the generated function is simply called `parse`.

3.3.3 The `%import` **directive**

The `%import` directive is used to include one grammar inside another. The string given in the directive should hold the path to a grammar file, and \ characters must be escaped. By default, all declarations appearing in the specified file are included in the resulting grammar, except for `%start`, `%entry`, and `%name` declarations. However, individual tokens or nonterminals can be dropped by listing them in the `%dropping` clause of an `%import` declaration. Since nonterminal definitions are cumulative (Section 3.4), the imported nonterminals can be extended with new productions simply by listing them. The final grammar must, of course, ensure that all used tokens and nonterminals are defined.

3.3.4 The `%keywords` **directive**

When a syntax error is discovered, `ml-antlr` attempts to find a single-token repair to the input that will allow the parse to continue. Changes to the input involving keywords can drastically alter the meaning of the input, so it is usually desirable to favor non-keyword repairs. The `%keywords` directive is used to tell `ml-antlr` which tokens should be considered keywords.

3.3.5 The `%name` **directive**

The prefix to use for the name of the generated parser functor is specified using `%name`. In addition to the functor, `ml-antlr` will generate a module to define the `token` datatype. If the declaration

```
%name Example;
```

appears in the specification, then the parser functor will be named `ExampleParseFn` and the tokens module will be called `ExampleTokens`.

3.3.6   The `%header` **directive**

The `%header` directive allows one to specify the parser's functor. This directive is useful when you want to add additional parameters to the functor, but the declaration must include the `Lex` structure with signature `ANTLR_LEXER`. For example,

```
%header (
    functor ExampleParseFn (
        structure Lex : ANTLR_LEXER
        structure Extras : EXTRA_SIG));
```

Note: the `%header` directive was added in SML/NJ version 110.72.

3.3.7   The `%refcell` **directive**

Because semantic actions must be pure (for backtracking and error repair), they cannot make use of standard reference cells to communicate information. Nonterminals may inherit attributes (Section 3.4.2), which allows information to flow downward, but in some cases flowing information this way can become extremely tedious. For example, a data structure may only need to be updated at a single point in the grammar, but in order to properly thread this state through the grammar, an inherited attribute would have to be added and propagated through every nonterminal.

The `%refcell` directive is used to declare a backtracking-safe reference cell and make it available to all semantic actions. Reference cells are declared by giving the name, type, and initial value for the cell. Each cell is bound in the semantic actions as a standard SML `ref` value. Thus, for example, we might have the following specification fragment:

```
%refcell symbols : StringSet.set = ( StringSet.empty );

exp
  : INT
  | (exp)
  | ID => ( symbols := StringSet.add(!symbols, ID); ID )
  ;
```

The result of this fragment is that all symbol uses are tracked, in any use of the `exp` nonterminal, but without having to manually thread the data structure state through the grammar.

The final contents of a reference cell is returned as an extra result from the parse following the repair list.

3.3.8   The `%start` **directive**

A particular nonterminal must be designated as the start symbol for the grammar. The start symbol can be specified using `%start`; otherwise, the first nonterminal defined is assumed to be the start symbol.

3.3.9   The %tokens directive

The alphabet of the parser is defined using %tokens. The syntax for this directive resembles a datatype declaration in SML, except that optional abbreviations for tokens may be defined. For example:

```
%tokens
  : KW_let  ("let")  | KW_in   ("in")
  | ID of string     | NUM of Int.int
  | EQ      ("=")    | PLUS    ("+")
  | LP      ("(")    | RP      (")")
  ;
```

Within nonterminal definitions, tokens may be referenced either by their name or abbreviation; the latter must always be double-quoted.

3.4   Nonterminal definitions

The syntax of nontermal definitions is given in Figure 3.1. As an illustration of the grammar, consider the following example, which defines a nonterminal with three productions, taking a formal parameter env:

```
atomicExp(env)
  : ID => ( valOf(AtomMap.find (env, Atom.atom ID)) )
  | NUM
  | "(" exp@(env) ")"
  ;
```

Note that actions are only allowed at the end of a production, and that they are optional.

As with most directives, the non-terminal definitions are cumulative. For example, the definition of atomicExp above could also be written as three separate rules.

```
atomicExp(env) : ID => ( valOf(AtomMap.find (env, Atom.atom ID)) );
atomicExp(env) : NUM;
atomicExp(env) : "(" exp@(env) ")";
```

3.4.1   Extended BNF constructions

In standard BNF syntax, the right side of a production is a simple string of symbols. Extended BNF allows regular expression-like operators to be used: *, +, and ? can follow a symbol, denoting 0 or more, 1 or more, or 0 or 1 occurrences respectively. In addition, parentheses can be used within a production to enclose a *subrule*, which may list several |-separated alternatives, each of which may have its own action. In the following example, the nonterminal item_list matches a semicolon-terminated list of identifiers and integers:

```
item_list : (( ID | INT ) ";")* ;
```

All of the extended BNF constructions have implications for the actions of a production; see Section 3.4.5 for details.

3.4.2   Inherited attributes

In most parsers, information can flow upward during the parse through actions, but not downward. In attribute grammar terminology, the former refers to *synthesized* attributes, while the latter refers to *inherited attributes*. Since `ml-antlr` is a predictive parser, it allows both kinds of attributes. Inherited attributes are treated as parameters to nonterminals, which can be used in their actions or semantic predicates. Formal parameters are introduced by enclosing them in parentheses after the name of a nonterminal and before its production list; the list of parameters will become a tuple. In the following, the nonterminal `expr` takes a single parameter called `env`:

```
expr(env) : (* ... *) ;
```

If a nonterminal has a formal parameter, any use of that nonterminal is required to apply it to an actual parameter. Actual parameters are introduced in a production by giving the name of a nonterminal, followed by the @ sign, followed by the code to compute the parameter. For example:

```
assignment : ID ":=" expr@(Env.emptyEnv) ;
```

3.4.3   Selective backtracking

Sometimes it is inconvenient or impossible to construct a nonterminal definition which can be unambiguously resolved with finite lookahead.The `%try` keyword can be used to mark ambiguous *productions* for selective backtracking. For backtracking to take place, each involved production must be so marked. Consider the following:

```
A : %try B* ";"
  | %try B* "(" C+ ")"
  ;
```

As written, the two productions cannot be distinguished with finite lookahead, since they share an arbitrary long prefix of `B` nonterminal symbols. Adding the `%try` markers tells `ml-antlr` to attempt to parse the first alternative, and if that fails to try the second. Another way to resolve the ambiguity is the use of subrules, which do not incur a performance penalty:

```
A : B* ( ";"
       | "(" C+ ")"
       )
    ;
```

This is essentially *left-factoring*. See Section 3.5 for more guidance on working with the $LL(k)$ restriction.

3.4.4 Semantic predicates

A production can be qualified by a *semantic predicate* by introducing a `%where` clause. Even if the production is syntactically matched by the input, it will not be used unless its semantic predicate evaluates to `true`. A `%where` clause can thus introduce context-sensitivity into a grammar. The following example uses an inherited `env` attribute, containing a variable-value environment:

```
atomicExp(env)
  : ID %where ( AtomMap.inDomain(env, Atom.atom ID) )
        => ( valOf(AtomMap.find (env, Atom.atom ID)) )
  | NUM
  | "(" exp@(env) ")"
  ;
```

In this example, if a variable is mentioned that has not been defined, the error is detected and reported during the parse as a syntax error.

Semantic predicates are most powerful when combined with selective backtracking. The combination allows two syntactically identical phrases to be distinguished by contextual, semantic information.

3.4.5 Actions

Actions for productions are just SML code enclosed in parentheses. Because of potential backtracking and error repair, action code should be pure (except that they may update `ml-antlr refcells`; see the `%refcell` directive).

In scope for an action are all the user definitions from the `%defs` directive. In addition, the formal parameters of the production are in scope, as are the semantic yield of all symbols to the left of the action (the yield of a token is the data associated with that token's constructor). In the following example, the first action has `env` and `exp` in scope, while the second action has `env` and `NUM` in scope:

```
atomicExp(env)
  : "(" exp@(env) ")" => ( exp )
  | NUM => ( NUM )
  ;
```

Notice also that the actual parameter to `exp` in the first production is `env`, which is in scope at the point the parameter is given; `exp` itself would not be in scope at that point.

An important aspect of actions is naming: in the above example, `exp` and `NUM` were the default names given to the symbols in the production. In general, the default name of a symbol is just the symbol's name. If the same name appears multiple times in a production, a number is appended to the name of each yield, start from 1, going from left to right. A subrule (any items enclosed in parentheses) is by default called `SR`. Any default name may be overriden using the syntax `name=symbol`. Overriding a default name does *not* change the automatic number for other default names. Consider:

```
foo : A bar=A A ("," A)* A*
    ;
```

In this production, the names in scope from left to right are: A1, bar, A3, SR, A4.

The EBNF operators *, + and ? have a special effect on the semantic yield of the symbols to which they are applied. Both * and + yield a *list* of the type of their symbol, while ? yields an option. For example, if ID* appeared in a production, its default name would be ID, and if the type of value of ID was string, it would yield a string list; likewise ID? would yield a string option.

Subrules can have embedded actions that determine their yield:

```
plusList : ((exp "+" exp => ( exp1 + exp2 )) ";" => ( SR ))* => ( SR )
```

The plusList nonterminal matches a list of semicolon-terminated additions. The innermost subrule, containing the addition, yields the value of the addition; that subrule is contained in a larger subrule terminated by a semicolon, which yield the value of the inner subrule. Finally, the semicolon-terminated subrule is itself within a subrule, which is repeated zero or more times. Note that the numbering scheme for names is restarted within each subrule.

Actions are *optional*: if an action is not specified, the default behavior is to return all nonterminals and non-nullary tokens in scope. Thus, the last example can be written as

```
plusList : ((exp "+" exp => ( exp1 + exp2 )) ";")*
```

since "+" and ";" represent nullary token values.

3.5 The $LL(k)$ restriction

When working with any parser, one must be aware of the restrictions is algorithm places on grammars. When ml-antlr analyzes a grammar, it attempts to create a prediction-decision tree for each nonterminal. In the usual case, this decision is made using lookahead token sets. The tool will start with $k = 1$ lookahead and increment up to a set maximum until it can uniquely predict each production. Subtrees of the decision tree remember the tokens chosen by their parents, and take this into account when computing lookahead. For example, suppose we have two productions at the top level that generate the following sentences:

```
prod1 ==> AA
prod1 ==> AB
prod1 ==> BC
prod2 ==> AC
prod2 ==> C
```

At $k = 1$, the productions can generate the following sets:

```
prod1 {A, B}
prod2 {A, C}
```

and $k = 2$,

```
prod1 {A, B, C}
prod2 {C, <EOF>}
```

Examining the lookahead sets alone, this grammar fragment looks ambiguous even for $k = 2$. However, `ml-antlr` will generate the following decision tree:

```
if LA(0) = A then
  if LA(1) = A or LA(1) = B then
    predict prod1
  else if LA(1) = C then
    predict prod2
else if LA(0) = B then
  predict prod1
else if LA(1) = C then
  predict prod2
```

   In `ml-antlr`, only a small amount of lookahead is used by default ($k = 3$). Thus, the following grammar is ambiguous for `ml-antlr`:

```
foo : A A A B
    | A A A A
    ;
```

and will generate the following error message:

```
Error: lookahead computation failed for 'foo',
with a conflict for the following productions:
  foo ::= A A A A EOF
  foo ::= A A A B EOF
The conflicting token sets are:
  k = 1: {A}
  k = 2: {A}
  k = 3: {A}
```

Whenever a lookahead ambiguity is detected, an error message of this form is given. The listed productions are the point of conflict. The `k = ...` sets together give examples that can cause the ambiguity, in this case an input of `AAA`.

   The problem with this example is that the two `foo` productions can only be distinguished by a token at $k = 4$ depth. This situation can usually be resolved using *left-factoring*, which lifts the common prefix of multiple productions into a single production, and then distinguishes the old productions through a subrule:

```
foo : A A A (B | A)
    ;
```

Recall that subrule alternatives can have their own actions:

```
    foo : A A A ( B => ( "got a B" )
                | A => ( "got an A" )
                )
        ;
```

making left-factoring a fairly versatile technique.

Another limitation of predictive parsing is *left-recursion*, where a nonterminal recurs without any intermediate symbols:

```
    foo : foo A A
        | B
        ;
```

Left-recursion breaks predictive parsing, because it is impossible to make a prediction for a left-recursive production without already having a prediction in hand. Usually, this is quite easily resolved using EBNF operators, since left-recursion is most often used for specifying lists. Thus, the previous example can be rewritten as

```
    foo : B (A A)*
        ;
```

which is both more readable and more amenable to $LL(k)$ parsing.

3.6   Position tracking

`ml-antlr` includes built-in support for propagating position information. Because the lexer module is required to provide a `getPos` function, the tokens themselves do not need to carry explicit position information. A position *span* is a pair to two lexer positions (the type `AntlrStreamPos.span` is an abbreviation for `AntlrStreamPos.pos * AntlrStreamPos.pos`). Within action code, the position span of any symbol (token, non-terminal, subrule) is available as a value; if the yield of the symbol is named `Sym`, its span is called `Sym_SPAN`. Note that the span of a symbol after applying the * or + operators is the span of the entire matched list:

```
    foo : A* => (* A_SPAN starts at the first A and ends at the last *)
```

In addition, the span of the entire current production is available as `FULL_SPAN`.

3.7   Using the generated code

When `ml-antlr` is run, it generates a tokens module and a parser functor. If the parser is given the name Xyz via the `%name` directive, these structures will be called `XyzParseFn` and `XyzTokens` respectively. The tokens module will contain a single datatype, called `token`. The data constructors for the `token` type have the same name and type as those given in the `%tokens` directive; in addition, a nullary constructor called `EOF` will be available.

The generated parser functor includes the following:

```
val parse :
  (Lex.strm -> ParserToks.token * AntlrStreamPos.span * Lex.strm) ->
  Lex.strm ->
  result_ty option * strm * ParserToks.token AntlrRepair.repair list
```

where `result_ty` is the type of the semantic action for the grammar's start symbol. The `parse` function is given a lexer function and a stream. The result of a parse is the semantic yield of the parse, the value of the stream at the end of the parse, and a list of error repairs. If an unrepairable error occurred, `NONE` is returned for the yield of the parse.

Note that if the start symbol for the grammar includes an inherited attribute (or a tuple of attributes), it will appear as an additional, curried parameter to the parser following the lexer parameter. Suppose, for example, that a grammar has a start symbol with an inherited `Int.int AtomMap.map`, and that the grammar yields `Int.int` values. The type of its `parse` function is as follows:

```
val parse :
  (strm -> ParserToks.token * strm) ->
  Int.int AtomMap.map ->
  strm ->
  Int.int option * strm * ParserToks.token AntlrRepair.repair list
```

The `AntlrRepair` module is part of the `ml-lpt-lib` library; it is fully described in Chapter 4. It includes a function `repairToString`:

```
val repairToString :
  ('token -> string) -> AntlrStreamPos.sourcemap ->
  'token repair -> string
```

Likewise, the tokens module (`ParserTokens` in this example) includes a function:

```
val toString : token -> string
```

Thus, although error reporting is customizable, a reasonable default is provided, as illustrated below:

```
let
  val sm = AntlrStreamPos.mkSourcemap()
  val (result, strm', errs) = Parser.parse (Lexer.lex sm) strm
  val errStrings =
    map (AntlrRepair.repairToString ParserTokens.toString sm)
        errs
in
  print (String.concatWith "\n" errStrings)
end
```

The `toString` function will convert each token to its symbol as given in a `%tokens` directive, using abbreviations when they are available. By substituting a different function for `toString`, this behavior can be altered.

Chapter 4

The `ml-lpt-lib` library

4.1 The `AntlrStreamPos` structure

```
structure AntlrStreamPos : sig

  type pos = Position.int
  type span = pos * pos
  type sourceloc = { fileName : string option, lineNo : int, colNo : int }
  type sourcemap

  (* the result of moving forward an integer number of characters *)
  val forward : pos * int -> pos

  val mkSourcemap  : unit   -> sourcemap
  val mkSourcemap' : string -> sourcemap

  val same : sourcemap * sourcemap -> bool

  (* log a new line occurence *)
  val markNewLine : sourcemap -> pos -> unit
  (* resychronize to a full source location *)
  val resynch     : sourcemap -> pos * sourceloc -> unit

  val sourceLoc : sourcemap -> pos -> sourceloc
  val fileName  : sourcemap -> pos -> string option
  val lineNo    : sourcemap -> pos -> int
  val colNo     : sourcemap -> pos -> int
  val toString  : sourcemap -> pos -> string
  val spanToString : sourcemap -> span -> string

end
```

25

4.2   The `AntlrRepair` structure

```
structure AntlrRepair : sig

  datatype 'tok repair_action
    = Insert of 'tok list
    | Delete of 'tok list
    | Subst of {
        old : 'tok list,
        new : 'tok list
      }
    | FailureAt of 'tok

  type 'a repair = AntlrStreamPos.pos * 'tok repair_action

  val actionToString :
        ('tok -> string)
          -> 'tok repair_action -> string

  val repairToString :
        ('tok -> string)
          -> AntlrStreamPos.sourcemap
            -> 'tok repair -> string

end
```

Chapter 5

A complete example

This chapter gives a complete example of a simple calculator language implemented using both `ml-ulex` and `ml-antlr`. Figure 5.1 gives the CM file for the project.

```
Library

  structure CalcLexer
  functor CalcParseFn
  structure CalcTest

is
  $/basis.cm
  $/smlnj-lib.cm
  $/ml-lpt-lib.cm

  calc.grm : ml-antlr
  calc.lex : ml-ulex
  calc-test.sml
```

Figure 5.1: The CM file: `sources.cm`

```
%name CalcLexer;

%let digit = [0-9];
%let int = {digit}+;
%let alpha = [a-zA-Z];
%let id = {alpha}({alpha} | {digit})*;

%defs (
  structure T = CalcTokens
  type lex_result = T.token
  fun eof() = T.EOF
);

let      => ( T.KW_let );
in       => ( T.KW_in );
{id}     => ( T.ID yytext );
{int}    => ( T.NUM (valOf (Int.fromString yytext)) );
"="      => ( T.EQ );
"+"      => ( T.PLUS );
"-"      => ( T.MINUS );
"*"      => ( T.TIMES );
"("      => ( T.LP );
")"      => ( T.RP );
" " | \n | \t
         => ( continue() );
.        => ( (* handle error *) );
```

Figure 5.2: The ml-ulex specification: `calc.lex`

```
%name Calc;

%tokens
  : KW_let  ("let")  | KW_in   ("in")
  | ID of string     | NUM of Int.int
  | EQ       ("=")   | PLUS    ("+")
  | TIMES   ("*")    | MINUS   ("-")
  | LP       ("(")   | RP      (")")
  ;

exp(env)
  : "let" ID "=" exp@(env)
    "in" exp@(AtomMap.insert(env, Atom.atom ID, exp1))
      => ( exp2 )
  | addExp@(env)
  ;

addExp(env)
  : multExp@(env) ("+" multExp@(env))*
      => ( List.foldr op+ 0 (multExp::SR) )
  ;

multExp(env)
  : prefixExp@(env) ("*" prefixExp@(env))*
      => ( List.foldr op* 1 (prefixExp::SR) )
  ;

prefixExp(env)
  : atomicExp@(env)
  | "-" prefixExp@(env)
      => ( ~prefixExp )
  ;

atomicExp(env)
  : ID
      => ( valOf(AtomMap.find (env, Atom.atom ID)) )
  | NUM
  | "(" exp@(env) ")"
  ;
```

Figure 5.3: The ml-antlr specification: `calc.grm`

```
structure CalcTest =
  struct

    structure CP = CalcParseFn(CalcLexer)

  (* val calc : TextIO.instream -> Int.int *)
    fun calc instrm = let
      val sm = AntlrStreamPos.mkSourcemap()
      val lex = CalcLexer.lex sm
      val strm = CalcLexer.streamifyInstream instrm
      val (r, strm', errs) = CP.parse lex AtomMap.empty strm
      in
        print (String.concatWith "\n"
            (map (AntlrRepair.repairToString
                   CalcTokens.toString sm)
                 errs));
        r
      end

  end
```

Figure 5.4: The driver: `calc-test.sml`